

Proximal Methods in Numerical Optimization

Lecture III – Proximal Gradient Methods

Alberto DE MARCHI

University of the Bundeswehr Munich

`alberto.demarchi@unibw.de`

`aldma.github.io`

Sapienza University of Rome – May 6, 2026



these slides are under development: please email me for corrections and suggestions



Outline

Settings

Proximal gradient methods

Convergence theory

Practical aspects

Toward (quasi-)Newton methods

Structured Optimization Problem

Problem (P)

$$\underset{x \in \mathbb{R}^n}{\text{minimize}} \quad \varphi(x) := f(x) + g(x)$$

Smooth part f

- ▶ $f : \mathbb{R}^n \rightarrow \mathbb{R}$ differentiable
- ▶ ∇f globally/locally Lipschitz/continuous
- ▶ possibly nonconvex

Nonsmooth part g

- ▶ $g : \mathbb{R}^n \rightarrow \overline{\mathbb{R}}$ proper, lsc (unif. prox-bounded)
- ▶ prox-friendly: “simple”, often convex
- ▶ $\text{dom } g$ encodes constraints

Motivating Examples

Example 1: LASSO

$$\underset{x \in \mathbb{R}^n}{\text{minimize}} \quad \frac{1}{2} \|Ax - b\|^2 + \lambda \|x\|_1$$

- ▶ f is $\mathcal{C}^\infty$, convex,
 $\nabla f(x) = A^\top (Ax - b)$
- ▶ g promotes sparsity
- ▶ soft-thresholding $\text{prox}_{\alpha g}$

Example 2: Constrained QP

$$\underset{x \in C}{\text{minimize}} \quad \frac{1}{2} x^\top Qx + q^\top x$$

- ▶ f is $\mathcal{C}^\infty$, quadratic
- ▶ $g = \delta_C$ (indicator of closed set C)
- ▶ $\text{prox}_{\alpha g} = \text{proj}_C$
- ▶ recovers projected gradient method

Further examples: group LASSO, nuclear norm minimization, sparse inverse covariance, neural-network training with ℓ_1 regularization, phase retrieval, matrix factorization...

Gradient Descent: A Quick Review

Problem: minimize _{x} $f(x)$ with $f \in \mathcal{C}^1$.

Gradient Descent (GD)

$$x_{k+1} = x_k - \alpha_k \nabla f(x_k)$$

Three views of the same step:

1. **Steepest descent:** follow direction $-\nabla f(x_k)$, scaled by α_k
2. **Linearization:** $x_{k+1} \in \arg \min_x [f(x_k) + \langle \nabla f(x_k), x - x_k \rangle] \dots$ (unbounded, TR)
3. **Quadratic model:**

$$Q(x; x_k, \alpha) := f(x_k) + \langle \nabla f(x_k), x - x_k \rangle + \frac{1}{2\alpha} \|x - x_k\|^2$$

$$\alpha_k > 0 \quad \implies \quad x_{k+1} \in \arg \min_x Q(x; x_k, \alpha_k) = x_k - \alpha_k \nabla f(x_k)$$

The quadratic model is the key to generalizing to structured problems $f + g$.


Proximal gradient methods

Quadratic Upper Models: Toward Majorization-Minimization

Suppose ∇f is L -Lipschitz, namely

$$f(y) \leq f(x) + \langle \nabla f(x), y - x \rangle + \frac{L}{2} \|y - x\|^2 = Q(y; x, 1/L)$$

The quadratic model $Q(\cdot; x, \alpha)$ is a (local) **upper bound** on f for all $\alpha \in (0, 1/L)$.



Minimizing the majorization model $Q(\cdot; x_k, 1/L) + g$:

$$\begin{aligned} x_{k+1} &\in \arg \min_x \{ Q(x; x_k, 1/L) + g(x) \} \\ &= \arg \min_x \left\{ f(x_k) + \langle \nabla f(x_k), x - x_k \rangle + \frac{L}{2} \|x - x_k\|^2 + g(x) \right\} \\ &= \arg \min_x \left\{ g(x) + \frac{L}{2} \left\| x - x_k + \frac{\nabla f(x_k)}{L} \right\|^2 \right\} = \text{prox}_{\alpha g} (x_k - \alpha \nabla f(x_k)) \end{aligned}$$

where $\alpha \equiv 1/L$. This is the **proximal gradient step** for $f + g$ at x_k with stepsize α .

Proximal Gradient Method (PGM)

Algorithm: basic PGM

```
1: Choose  $x^0 \in \text{dom } g$ , stepsize  $\alpha > 0$ 
2: for  $k = 0, 1, 2, \dots$  do
3:    $x_{k+1} \in \text{prox}_{\alpha g}(x_k - \alpha \nabla f(x_k))$ 
4: end for
```

Interpretation:

- ▶ Half-step: *gradient* on f
- ▶ Half-step: *proximal* on g
- ▶ Operator splitting: forward-backward, explicit-implicit

Fixed-point iteration.

Characterization for **convex** g :

$$\begin{aligned} x^* \text{ stationary} &\iff 0 \in \partial g(x^*) + \nabla f(x^*) \\ &\iff x^* = \text{prox}_{\alpha g}(x^* - \alpha \nabla f(x^*)) \end{aligned}$$

Projected Gradient as a Special Case

Let $g = \delta_C$ (indicator of closed set $C \subseteq \mathbb{R}^n$):

$$\delta_C(x) = \begin{cases} 0 & x \in C \\ +\infty & x \notin C \end{cases}$$

Then:

$$\text{prox}_{\alpha \delta_C}(z) = \arg \min_{u \in C} \|u - z\|^2 = \text{proj}_C(z)$$

Projected Gradient Descent (PGD)

$$x_{k+1} \in \text{proj}_C(x_k - \alpha_k \nabla f(x_k))$$

Examples: box constraints (ℓ_∞ ball), simplex projection, nuclear norm ball, PSD cone, complementarity constraints...

Block Coordinate Proximal Gradient

When $x = (x_1, \dots, x_B)$ is partitioned into blocks:

$$\varphi(x) := f(x) + \sum_{b=1}^B g_b(x_b)$$

Block Coordinate Descent (BCD-PGM)

- 1: **for** $k = 0, 1, \dots$ **do**
- 2: Pick block $b_k \in \{1, \dots, B\}$
- 3: $(x_{b_k})_{k+1} = \text{prox}_{\alpha g_{b_k}}((x_{b_k})_k - \alpha \nabla_{b_k} f(x_k))$
- 4: $(x_{\neg b_k})_{k+1} = (x_{\neg b_k})_k$
- 5: **end for**

- ▶ Cheap per-iteration cost (partial gradient)
- ▶ Cyclic or random selection of blocks b_k for convergence to stationary points
- ▶ Suited for matrix factorization, dictionary learning, sparse NMF

Convergence theory

Global Lipschitz Gradient Assumption

Assumption (L)

∇f is L -Lipschitz continuous: there exists a finite L such that

$$\|\nabla f(x) - \nabla f(y)\| \leq L \|x - y\| \quad \forall x, y \in \mathbb{R}^n.$$

► **Descent lemma:** $f(y) \leq f(x) + \langle \nabla f(x), y - x \rangle + \frac{L}{2} \|y - x\|^2$ for all x, y .

$$f(x_+) \leq f(x) + \langle \nabla f(x), x_+ - x \rangle + \frac{L}{2} \|x_+ - x\|^2$$

► **Prox-grad step:** $x_+ \in \text{prox}_{\alpha g}(x - \alpha \nabla f(x)) = \arg \min_x \left\{ g(z) + \frac{1}{2\alpha} \|z - x\|^2 \right\}.$

$$g(x_+) + \frac{1}{2\alpha} \|x_+ - x + \alpha \nabla f(x)\|^2 \leq g(x) + \frac{\alpha}{2} \|\nabla f(x)\|^2$$

$$\implies \varphi(x_+) \leq \varphi(x) + \frac{\alpha L - 1}{2\alpha} \|x_+ - x\|^2$$

Global Lipschitz Gradient Assumption II

Assumption (L)

∇f is L -Lipschitz continuous: there exists a finite L such that

$$\|\nabla f(x) - \nabla f(y)\| \leq L \|x - y\| \quad \forall x, y \in \mathbb{R}^n.$$

$$\varphi(x_+) \leq \varphi(x) + \frac{\alpha L - 1}{2\alpha} \|x_+ - x\|^2$$

With $\alpha = \sigma/L$, $\sigma \in (0, 1)$, there is **sufficient decrease**:

$$x_{k+1} \in \text{prox}_{\alpha g}(x_k - \alpha \nabla f(x_k)) \implies \varphi(x_{k+1}) \leq \varphi(x_k) - \frac{1 - \sigma}{2\alpha} \|x_{k+1} - x_k\|^2$$

Issue: L often unknown or overly conservative $\Rightarrow$ use backtracking to find good α .

Gradient Mapping

Gradient mapping

Generalized gradient for $f + g$:

$$G_\alpha(x) := \frac{x - \text{prox}_{\alpha g}(x - \alpha \nabla f(x))}{\alpha} = \frac{x - x_+}{\alpha}$$

Why it matters:

- ▶ Reduces to $\nabla f(x)$ when $g = 0$ (unconstrained, smooth).
- ▶ **Stationarity:** $G_\alpha(x^*) = 0 \iff 0 \in \nabla f(x^*) + \partial g(x^*)$ for **convex** g .
- ▶ $\|G_\alpha(x)\|$ is a natural *measure of stationarity* at x .

$$\varphi(x_+) \leq \varphi(x) - \alpha \frac{1 - \sigma}{2} \|G_\alpha(x)\|^2 = \varphi(x) - \frac{1 - \sigma}{2\alpha} \|x_+ - x\|^2$$

PGM step: $x_{k+1} = x_k - \alpha G_\alpha(x_k)$.

- ▶ PGM as GD with *gradient mapping*.
- ▶ Convergence requires $\|G_\alpha(x^k)\| \rightarrow 0$.

Sufficient Decrease

Lemma (Sufficient Decrease).

Under Assumption (L), with $\alpha \in (0, 1/L)$ and some $\sigma \in (0, 1)$:

$$\varphi(x_{k+1}) \leq \varphi(x_k) - \frac{\sigma\alpha}{2} \|G_\alpha(x_k)\|^2 \leq \varphi(x_k)$$

If $\varphi := f + g$ is bounded from below by $\varphi_\star$, then starting with $x_0 \in \text{dom } \varphi$ and telescoping

$$\infty > \varphi(x_0) - \varphi_\star \geq \sum_{k=0}^K \varphi(x_k) - \varphi(x_{k+1}) \geq \frac{\sigma\alpha}{2} \sum_{k=0}^K \|G_\alpha(x_k)\|^2$$

meaning that $\|G_\alpha(x_k)\| \rightarrow 0$.

Backtracking Line Search

When L is unknown or large, use **adaptive stepsize**, typically based on **backtracking**:

Backtracking: one PGM iteration

- 1: **Input:** $x_k, \bar{\alpha} > 0, \beta \in (0, 1)$
- 2: $\alpha_k \leftarrow \bar{\alpha}$
- 3: $x_{k+1} \leftarrow \text{prox}_{\alpha_k g}(x_k - \alpha_k \nabla f(x_k))$
- 4: **while** $f(x_{k+1}) > Q(x_{k+1}; x_k, \alpha_k)$ **do**
- 5: $\alpha_k \leftarrow \beta \alpha_k$
- 6: $x_{k+1} \leftarrow \text{prox}_{\alpha_k g}(x_k - \alpha_k \nabla f(x_k))$
- 7: **end while**

where $Q(y; x, \alpha) := f(x) + \langle \nabla f(x), y - x \rangle + \frac{1}{2\alpha} \|y - x\|^2$ is a quadratic model of f around x .

- ▶ Adaptive to local curvature
- ▶ Guarantees sufficient decrease at each iteration
- ▶ Terminates in finite steps (since $\alpha \leq 1/L$ satisfies the condition)
- ▶ ... even for locally Lipschitz ∇f , or less [DMT22, KM22]

Convergence: Convex Case

Additional Assumption (C):

f and g are convex; $\varphi_{\star} = \min \varphi > -\infty$.

Theorem (Sublinear Rate, $\mathcal{O}(1/k)$).

Under (L) + (C), with fixed $\alpha = 1/L$:

$$\varphi(x_k) - \varphi_{\star} \leq \frac{L \|x_0 - x_{\star}\|^2}{2k}$$

Key steps in proof:

- ▶ Sufficient decrease + convexity of $g \implies$ one-step progress inequality.
- ▶ Telescope the sufficient-decrease inequality over k steps.
- ▶ Use convexity of φ to bound $\varphi(x_k)$ by $\frac{1}{k} \sum_j \varphi(x_j)$.

Linear Convergence Under Strong Convexity

Assumption (SC):

f is μ -strongly convex ($\mu > 0$).

Theorem (Linear Rate).

Under (L) + (SC), with $\alpha = 1/L$:

$$\|x_k - x_\star\|^2 \leq \left(1 - \frac{\mu}{L}\right)^k \|x_0 - x_\star\|^2$$

Condition number $\kappa = L/\mu$ governs the rate.

Derivation: The key contraction is

$$\|\mathcal{T}_\alpha(x) - \mathcal{T}_\alpha(y)\|^2 \leq \left(1 - \frac{\mu}{L}\right) \|x - y\|^2$$

where $\mathcal{T}_\alpha(x) := \text{prox}_{\alpha g}(x - \alpha \nabla f(x))$.

Convergence: Nonconvex Case

Without convexity, we cannot expect $\varphi(x_k) \rightarrow \inf \varphi$. Instead:

Theorem (Stationarity in $\mathcal{O}(K)$).

Under Assumption (L), backtracking PGM produces:

$$\min_{0 \leq k \leq K} \|G_{\alpha}(x_k)\|^2 \leq \frac{2}{\alpha_{\min} K} (\varphi(x_0) - \varphi_{\star})$$

Hence $G_{\alpha}(x_k) \rightarrow 0$ along a subsequence.

Proof sketch:

- ▶ Sufficient decrease: $\varphi(x_{k+1}) \leq \varphi(x_k) - \frac{\alpha_k}{2} \|G_{\alpha_k}(x_k)\|^2$.
- ▶ Telescope: $\sum_{k=0}^{K-1} \frac{\alpha_k}{2} \|G_{\alpha_k}(x_k)\|^2 \leq \varphi(x_0) - \varphi_{\star} < \infty$.
- ▶ $\alpha_k \geq \alpha_{\min} := \beta/L > 0$.

If ∇f not globally Lipschitz, $\alpha_{\min}$ depends on geometry around accumulation points...

Kurdyka–Łojasiewicz (KL) Framework

For finer convergence results in the nonconvex case, we use:

KL Property.

Function φ satisfies KL at $x_\star$ if there exists a desingularizing function ψ such that

$$\psi'(\varphi(x) - \varphi(x_\star)) \|\partial\varphi(x)\| \geq 1$$

holds for all x near $x_\star$.

KL covers:

- ▶ Real-analytic, semi-algebraic functions (polynomials, ℓ_1 -regularized, ...)
- ▶ Piecewise smooth functions
- ▶ Practically all functions encountered in ML/stats

Convergence with KL: whole sequence

Under KL property and sufficient decrease, PGM sequence **converges** to a critical point with rate determined by ψ (linear/sublinear depending on KL exponent).

Accelerated Gradient Methods: Nesterov's Idea

Can we do better than $\mathcal{O}(1/k)$ for convex problems?

Nesterov (1983): Yes! Using **momentum**.

Key idea — extrapolation: instead of applying prox-grad at x_k , apply it at a *momentum point / looking ahead*:

$$y_k = x_k + \tau_k(x_k - x_{k-1})$$

where $\{\tau_k\}$ is a carefully chosen sequence.

Lower bound (Nesterov, 1983)

No first-order method can achieve better than $\mathcal{O}(1/k^2)$ in general.

FISTA's revolution: Fast ISTA / PGM

Algorithm: FISTA — Beck & Teboulle [BT09]

```
1:  $x_0 = y_1 \in \text{dom } g$ ,  $t_1 = 1$ ,  $\alpha = 1/L$   
2: for  $k = 1, 2, \dots$  do  
3:    $x_k = \text{prox}_{\alpha g}(y_k - \alpha \nabla f(y_k))$   
4:    $t_{k+1} = \frac{1}{2}(1 + \sqrt{1 + 4t_k^2})$   
5:    $y_{k+1} = x_k + \frac{t_k - 1}{t_{k+1}}(x_k - x_{k-1})$   
6: end for
```

Theorem.

Under (L)+(C), FISTA achieves $\mathcal{O}(1/k^2)$:

$$\varphi(x_k) - \varphi_* \leq \frac{2L \|x_0 - x_*\|^2}{(k+1)^2}.$$

Momentum factor $\tau_k \equiv \frac{t_k - 1}{t_{k+1}} \rightarrow 1^-$ asymptotically, $y_k \approx 2x_k - x_{k-1}$.

- ▶ Momentum can hurt near stationary points $\rightsquigarrow$ restarting (empirically much faster, especially for ill-conditioned problems)
- ▶ Monotone variants to ensure $\varphi(x^k)$ decreasing (slightly slower in practice)

Acceleration in the Nonconvex Setting

Warning

Nesterov momentum does *not* guarantee acceleration for nonconvex φ .

iPiano [OCBP14]: inertial PGM for nonconvex f and convex g :

$$x_{k+1} = \text{prox}_{\alpha_k g}(x_k - \alpha_k \nabla f(x_k) + \beta_k(x_k - x_{k-1}))$$

Convergence under KL property + sufficient decrease condition on (α_k, β_k) .

APGM-NC [LL15]: Nesterov acceleration with monitoring, for nonconvex f and nonconvex g . Achieves global convergence to stationary points in general, and complexity $\mathcal{O}(1/k^2)$ for convex problems.

Practical aspects

Termination Conditions

When do we stop? We need a computable, meaningful criterion.

Common criterion:

gradient mapping residual

$$\left\| G_{\alpha}(x^k) \right\| = \frac{1}{\alpha} \left\| x^k - x^{k+1} \right\| \leq \varepsilon$$

Additional/complementary conditions:

- ▶ Iterate change: $\|x^k - x^{k-1}\| \leq \varepsilon \max(1, \|x^k\|)$
- ▶ Objective change: $|\varphi(x^k) - \varphi(x^{k-1})| \leq \varepsilon \max(1, |\varphi(x^k)|)$
- ▶ Dual residual (g convex): $\|x^k - \text{prox}_g(x^k - \nabla f(x^k))\| \leq \varepsilon$
- ▶ Max iterations: $k \geq k_{\max}$

Practice: combine stationarity measure with iterate/objective change for robustness.

Termination Conditions (to be precise)

When do we stop? We need a computable **measure of optimality/stationarity**.

From the PGM update we have

$$\begin{aligned} 0 &\in \partial g(x_{k+1}) + \frac{x_{k+1} - x_k}{\alpha_k} + \nabla f(x_k) \\ \iff \quad \nabla f(x_{k+1}) - \nabla f(x_k) - \frac{x_{k+1} - x_k}{\alpha_k} &\in \partial g(x_{k+1}) + \nabla f(x_{k+1}) = \partial \varphi(x_{k+1}) \end{aligned}$$

Stationarity criterion:

$$\text{dist}(0, \partial \varphi(x_{k+1})) \stackrel{\text{PGM}}{\leq} \left\| \nabla f(x_{k+1}) - \nabla f(x_k) - \frac{x_{k+1} - x_k}{\alpha_k} \right\| \leq \varepsilon$$

Relaxing the Global Lipschitz Assumption

Global Lipschitz is often too strong:

- ▶ Unknown L , or L is very large (loose bound)
- ▶ Barrier / Nonconvex f : gradient may not be globally Lipschitz

Relaxations:

1. **Local Lipschitz + backtracking:** backtracking finds local L_k , guaranteed to terminate.
2. **Local descent lemma:** only require the backtracking condition (no global L).
3. **Proximal smoothness:** f prox-regular; weaker than convexity.
4. **Relative smoothness** [BBT17]: replace $\frac{1}{2} \|y - x\|^2$ by $D_h(y, x)$ (Bregman divergence); leads to *mirror descent* / *Bregman proximal gradient*.

Backtracking regularization

PGM with backtracking:

- ▶ prox-grad step: compute $x_+ \in \text{prox}_{\alpha g}(x - \alpha \nabla f(x))$
- ▶ sufficient decrease: check if $\varphi(x_+) \leq \varphi(x_k) - \frac{\sigma}{\alpha} \|x_+ - x\|^2$
- ▶ backtracking: update α or break

For **nonconvex** f , backtracking is the standard practical approach. It is an adaptive (quadratic) regularization, **not** a linesearch procedure.

For **convex** f and g , adaptive schemes without checking descent [MM20, MM24, LTSP25]

Nonmonotone Descent

Observation: Forcing $\varphi(x_{k+1}) \leq \varphi(x_k)$ every step can be overly restrictive

Key idea: replace *monotone* sufficient decrease condition with

$$\varphi(x_{k+1}) \leq \Phi_k - \frac{\sigma}{\alpha_k} \|x_{k+1} - x_k\|^2$$

where $\Phi_k \geq \varphi(x_k)$ and $\{\Phi_k\}$ ensures convergence.

Max-type — Grippo, Lampariello & Lucidi [GLL86]

Given $M \in \mathbb{N}$ and $m(k) := \min\{M, k\}$,

$$\Phi_k := \max_{0 \leq j \leq m(k)} \varphi(x_{k-j})$$

Mean-type — Zhang & Hager [ZH04], [DM23]

Given $\eta \in (0, 1]$,

$$\Phi_k := (1-\eta)\Phi_{k-1} + \eta\varphi(x_k), \quad \Phi_0 := \varphi(x_0)$$

Benefits:

- ▶ Accepts larger steps; avoids zigzagging
- ▶ Bounded nonmonotonicity; specializes to monotone
- ▶ Convergence guaranteed, possibly under additional assumptions (max. . .) [KM22, DM23]

Spectral (Barzilai–Borwein) Stepsizes

Idea: estimate the local Lipschitz constant from secant information.

Let $s_k := x_k - x_{k-1}$, $y_k := \nabla f(x_k) - \nabla f(x_{k-1})$. Mimic a diagonal quasi-Newton step:

BB Stepsizes — Barzilai & Borwein [BB88]

$$\alpha_k^{\text{BB1}} = \frac{\langle s^k, s^k \rangle}{\langle s^k, y^k \rangle} \quad (\text{short step}), \quad \alpha_k^{\text{BB2}} = \frac{\langle s^k, y^k \rangle}{\langle y^k, y^k \rangle} \quad (\text{long step})$$

- ▶ Approximate the inverse of the average curvature
- ▶ In practice, must pair with **nonmonotone** globalization
- ▶ Dramatic **speedup** in practice (close to quasi-Newton)

Spectral Proximal Gradient Method

Algorithm: SPG [BMR00, JKMW23] / SpaRSA [WNF09]

- 1: $x_0, \alpha_0^{\text{BB}} > 0, 0 < \alpha_{\min} < \alpha_{\max}$
- 2: **for** $k = 0, 1, 2, \dots$ **do**
- 3: Starting from α_k^{BB} , by backtracking find α_k that delivers suff. decrease
- 4: $x_{k+1} \in \text{prox}_{\alpha_k g}(x_k - \alpha_k \nabla f(x_k))$
- 5: Compute $s_{k+1} := x_{k+1} - x_k, y_{k+1} := \nabla f(x_{k+1}) - \nabla f(x_k)$
- 6: $\alpha_{k+1}^{\text{BB}} \leftarrow \text{clip}(\alpha^{\text{BB}}(s_{k+1}, y_{k+1}), \alpha_{\min}, \alpha_{\max})$
- 7: **end for**

In practice, often faster than FISTA

Toward (quasi-)Newton methods

Limitations of First-Order Proximal Gradient

First-order methods: only use ∇f (gradient information).

Issues:

- ▶ Rate $\mathcal{O}(1/k)$ (or $\mathcal{O}(1/k^2)$ with momentum) — slow for ill-conditioned problems
- ▶ Strong convexity gives linear rate $\mathcal{O}((1 - \mu/L)^k)$ but $\mu/L \ll 1$ in practice

Newton-type methods: use curvature (Hessian or approximation).

- ▶ Replace $\frac{1}{2\alpha} \|x - x_k\|^2$ by $\frac{1}{2}(x - x_k)^\top H_k(x - x_k)$ in the model of f
- ▶ *Variable metric prox-gradient:*

$$x_{k+1} \in \text{prox}_g^{H_k}(x_k - H_k^{-1} \nabla f(x_k))$$

- ▶ Challenge: prox^{H_k} may not be easy to compute for arbitrary H_k

Proximal Newton-type methods for convex structured minimization [PB13, LSS14]

Variable Metric Proximal Operator

Definition.

For SPD matrix H , the H -proximal operator is:

$$\text{prox}_g^H(z) := \arg \min_u \left\{ g(u) + \frac{1}{2}(u - z)^\top H(u - z) \right\}$$

Key issue: for general H , the above may have no closed form even when prox_g does.

Exceptions (tractable prox_g^H):

- ▶ $g = \lambda \|\cdot\|_1$ and *diagonal* $H \implies$ scaled soft-threshold
- ▶ $g = \delta_C$ and *diagonal* $H \implies$ scaled projection
- ▶ g separable and H diagonal $\implies$ coordinate-wise prox

Full (quasi-)Newton requires more sophisticated approaches to be practical.

Idea: combine fast directions with safe PGM steps.

Look at the **inclusion** problem

$$\text{find } x \text{ such that } 0 \in \mathcal{T}_\alpha(x) := \frac{1}{\alpha}[x - \text{prox}_{\alpha g}(x - \alpha \nabla f(x))].$$

We seek a zero of the (set-valued) fixed-point residual $\mathcal{T}_\alpha$.

Newton-type scheme for this root-finding problem:

$$x_{k+1} = x_k - H_k \mathcal{T}_\alpha(x_k)$$

where H_k capture the curvature of $\mathcal{T}_\alpha$ around x_k .

$\rightsquigarrow H_k$ from **quasi-Newton** formulas

But, what about globalization?

Key observation.

The Moreau envelope is continuous.

PANOC algorithm:

1. Compute $x_{k+\frac{1}{2}} = \text{prox}_{\alpha g}(x_k - \alpha \nabla f(x_k))$ (PGM step as anchor)
2. Build quasi-Newton direction on $\varphi_{\alpha}^{\text{FB}}$ using L-BFGS
3. Linesearch to find x_{k+1} satisfying $\varphi_{\alpha}^{\text{FB}}(x_{k+1}) \leq \varphi_{\alpha}^{\text{FB}}(x_k) - \frac{\sigma}{\alpha} \|x_{k+1} - x_k\|^2$

Forward-Backward Envelope (FBE)

Definition [PB13].

The *forward-backward envelope* of $\varphi := f + g$ is:

$$\varphi_{\alpha}^{\text{FB}}(x) := \min_u \left\{ f(x) + \langle \nabla f(x), u - x \rangle + \frac{1}{2\alpha} \|u - x\|^2 + g(u) \right\}$$

Properties (for suff. small $\alpha > 0$):

- ▶ $\varphi_{\alpha}^{\text{FB}}(x) \leq \varphi(x)$ for all x
- ▶ $\varphi_{\alpha}^{\text{FB}}(x) = \varphi(x) \iff x$ is a fixed point / critical / stationary
- ▶ for convex g , gradient mapping $G_{\alpha}(x) = \alpha \nabla \varphi_{\alpha}^{\text{FB}}(x)$

PANOC minimizes $\varphi_{\alpha}^{\text{FB}}$ by combining fast smooth quasi-Newton extrapolation steps with safe prox-grad updates.

PANOC: Properties and Convergence [STSP17, TSP18]

Structure:

- ▶ One prox evaluation per iteration
- ▶ One gradient evaluation per iteration
- ▶ L-BFGS memory update (low cost)
- ▶ Falls back to PGM if quasi-Newton step rejected

Convergence:

- ▶ $\|G_\alpha(x^k)\| \rightarrow 0$ (nonconvex)
- ▶ Superlinear rate near non-degenerate solutions
- ▶ Formally: R -superlinear under second-order sufficient conditions

PANOC⁺: extension to locally Lipschitz ∇f , also with inexact prox evaluations [DMT22]

Applications: embedded optimal control, MPC, robotics. Subsolver in OpEn [SFP20], alpaqa [PSP22], Bazinga [DMJKM23].

Regularized Proximal Newton-type Method

Since [PB13, LSS14], form the **second-order structured model** at iterate x_k :

$$m_k(x) := f(x_k) + \langle \nabla f(x_k), x - x_k \rangle + \frac{1}{2}(x - x_k)^\top B_k(x - x_k) + g(x)$$

where $B_k \approx \nabla^2 f(x_k)$ (exact Hessian or approximation) with regularization.

Proximal Newton-type subproblem

Solve

$$x_{k+1} \in \arg \min_x m_k(x)$$

Newton-type subproblem hardly available in closed-form, but a structured minimization: call PGM to solve it!

Ideas for practical proximal Newton-type methods: RPQN

RPQN [BFO19, Lec22]

Limited-memory quasi-Newton B_k with compact representation

$\implies$ small (memory M) system of equations to evaluate the scaled prox operator.

Nocedal's "Updating quasi-Newton matrices with limited storage" (1980)

Becker, Fadili & Ochs's "On quasi-Newton forward-backward splitting: Proximal calculus and convergence" (2019)

Kanzow & Lechner's "Efficient regularized proximal quasi-Newton methods for large-scale nonconvex composite optimization problems" (2024)

Ideas for practical proximal Newton-type methods: R2N

R2N [ABO22, DHO26]

- ▶ Solve the Newton-type proximal subproblem inexactly,
- ▶ ... obtaining at least the improvement of a prox-grad step (Cauchy decrease).
- ▶ Check if x_{k+1} delivers sufficient decrease for φ .
- ▶ If needed, recompute with an increased quadratic regularization.

Diouane, Habiboullah & Orban's "A proximal modified quasi-Newton method for nonsmooth regularized optimization" (2026)

Comparing the Methods: A Unified View

All methods minimize a *model* $m_k(x) \approx f(x) + g(x)$:

Method	Model $m_k(x)$	Per-iter cost
PGM	$\nabla f^\top d + \frac{1}{2\alpha} \ d\ ^2 + g$	1 grad, 1 prox
PGM + BB	same, $\alpha_k = \alpha^{\text{BB}}$, nonmono	1 grad, 1 prox
FISTA	same + momentum	1 grad, 1 prox
PANOC	envelope $\varphi_\alpha^{\text{FB}}$ + L-BFGS	1 grad, 1 prox, QN update
R2N	$\nabla f^\top d + \frac{1}{2} d^\top B_k d + g$	inexact inner PGM solve
RPQN	same with compact represent.	1 grad, 1 inner root-finder

Trade-off: richer model $\implies$ fewer outer iterations but more expensive subproblem.
The *right* method depends on the problem structure, cost of g , Hessian availability.

In practice, choose solver depending on most/least expensive oracles at hand.

Some comparisons in [DHO26].

Numerical examples

Sparse linear regression with ProximalAlgorithms.jl / RegularizedOptimization.jl

$$\underset{x}{\text{minimize}} \quad \underbrace{\frac{1}{2}\|Ax - b\|^2}_{=:f(x)} + \lambda R(x)$$

$$\nabla f(x) = A^\top (Ax - b)$$

$$\nabla f(x) - \nabla f(y) = A^\top A(x - y)$$

[$\rightsquigarrow$ tutorial: PG]

What about more interesting problems?

- ▶ **Proximal Newton Methods for Convex Composite Optimization**
Patrinos, Bemporad. CDC 2013.
- ▶ **A Simple and Efficient Algorithm for Nonlinear Model Predictive Control**
Stella, Themelis, Sopasakis, Patrinos. CDC 2017.
- ▶ **OpEn: Code Generation for Embedded Nonconvex Optimization**
Sopasakis, Fresk, Patrinos. IFAC 2020.

Exercise

Consider problems of the form

$$\underset{x \in \mathbb{R}^n}{\text{minimize}} \quad f(x) + g(x) + \delta_C(x)$$

where f is differentiable, $g \in \{\|\cdot\|_0, \|\cdot\|_1\}$ and $C := [\ell, u] \subseteq \mathbb{R}^n$ a nonempty set.

The code you write (in your favourite programming language) should take the following inputs: a function f , a flag to select g , vectors ℓ, u to define C , and a starting point x_0 . For 2-dimensional problem, you should also plot the constraint set C and the trajectory followed by the iterates $\{x_k\}$.

As concrete examples for f , you can take the Rosenbrock or Beale function from https://en.wikipedia.org/wiki/Test_functions_for_optimization.

Please send me a **single file with commented code before May 21th, 2026**.
Should you have any questions regarding the assignment, contact me asap 😊

Additional material, available online

- ▶ **Proximal Gradient Algorithms: Applications in Signal Processing.** Niccoló Antonello, Lorenzo Stella, Panos Patrinos, Toon van Waterschoot. EUSIPCO 2019 [https://www.eurasip.org/Proceedings/Eusipco/eusipco2019/wp-content/uploads/2019/10/T5_compressed.pdf] [<https://arxiv.org/pdf/1803.01621>]
- ▶ **Proximal Algorithms.** Neal Parikh, Stephen Boyd. Foundations and Trends in Optimization (2013) [https://web.stanford.edu/~boyd/papers/pdf/prox_algs.pdf]
- ▶ **Proximal Gradient Descent (and Acceleration).** Ryan Tibshirani. Convex Optimization 10-725 [<https://www.stat.cmu.edu/~ryantibs/convexopt/lectures/prox-grad.pdf>]
- ▶ **Proximal gradient method.** Lieven Vandenbergh. ECE236C (2022) [<https://www.seas.ucla.edu/~vandenbe/236C/lectures/proxgrad.pdf>]
- ▶ **Proximal gradient method.** Laurent El Ghaoui. EE227A (2010) [<https://people.eecs.berkeley.edu/~elghaoui/Teaching/EE227A/lecture18.pdf>]

References I

-  Aleksandr Y. Aravkin, Robert Baraldi, and Dominique Orban.
A proximal quasi-Newton trust-region method for nonsmooth regularized optimization.
SIAM Journal on Optimization, 32(2):900–929, 2022.
-  Jonathan Bartzilai and Jonathan M. Borwein.
Two-point step size gradient methods.
IMA Journal of Numerical Analysis, 8(1):141–148, 1988.
-  Heinz H. Bauschke, Jérôme Bolte, and Marc Teboulle.
A descent lemma beyond Lipschitz gradient continuity: First-order methods revisited and applications.
Mathematics of Operations Research, 42(2):330–348, 2017.
-  Stephen Becker, Jalal Fadili, and Peter Ochs.
On quasi-Newton forward-backward splitting: Proximal calculus and convergence.
SIAM Journal on Optimization, 29:2445–2481, 2019.
-  Ernesto G. Birgin, José Mario Martínez, and Marcos Raydan.
Nonmonotone spectral projected gradient methods on convex sets.
SIAM Journal on Optimization, 10(4):1196–1211, 2000.
-  Amir Beck and Marc Teboulle.
A fast iterative shrinkage-thresholding algorithm for linear inverse problems.
SIAM Journal on Imaging Sciences, 2(1):183–202, 2009.

References II

-  Youssef Diouane, Mohamed Laghdaf Habiboullah, and Dominique Orban.
A proximal modified quasi-newton method for nonsmooth regularized optimization.
SIAM Journal on Optimization, 2026.
-  Alberto De Marchi.
Proximal gradient methods beyond monotony.
Journal of Nonsmooth Analysis and Optimization, 4, 2023.
-  Alberto De Marchi, Xiaoxi Jia, Christian Kanzow, and Patrick Mehlitz.
Constrained composite optimization and augmented Lagrangian methods.
Mathematical Programming, 201(1):863–896, 9 2023.
-  Alberto De Marchi and Andreas Themelis.
Proximal gradient algorithms under local Lipschitz gradient continuity.
Journal of Optimization Theory and Applications, 194(3):771–794, 2022.
-  L. Grippo, F. Lampariello, and S. Lucidi.
A nonmonotone line search technique for Newton's method.
SIAM Journal on Numerical Analysis, 23(4):707–716, 1986.
-  Xiaoxi Jia, Christian Kanzow, Patrick Mehlitz, and Gerd Wachsmuth.
An augmented Lagrangian method for optimization problems with structured geometric constraints.
Mathematical Programming, 199(1):1365–1415, 2023.

References III



Christian Kanzow and Patrick Mehlitz.

Convergence properties of monotone and nonmonotone proximal gradient methods revisited.
Journal of Optimization Theory and Applications, 195(2):624–646, 2022.



Theresa Lechner.

Proximal Methods for Nonconvex Composite Optimization Problems.
PhD thesis, Julius-Maximilians-Universität Würzburg, 2022.



Huan Li and Zhouchen Lin.

Accelerated proximal gradient methods for nonconvex programming.
In C. Cortes, N. Lawrence, D. Lee, M. Sugiyama, and R. Garnett, editors, *Advances in Neural Information Processing Systems*, volume 28. Curran Associates, Inc., 2015.



Jason D. Lee, Yuekai Sun, and Michael A. Saunders.

Proximal Newton-type methods for minimizing composite functions.
SIAM Journal on Optimization, 24(3):1420–1443, 2014.



Puya Latafat, Andreas Themelis, Lorenzo Stella, and Panagiotis Patrinos.

Adaptive proximal algorithms for convex optimization under local lipschitz continuity of the gradient.
Mathematical Programming, 213(1):433–471, 2025.

References IV



Yura Malitsky and Konstantin Mishchenko.

Adaptive gradient descent without descent.

In Hal Daumé III and Aarti Singh, editors, *Proceedings of the 37th International Conference on Machine Learning*, volume 119 of *Proceedings of Machine Learning Research*, pages 6702–6712. PMLR, 13–18 Jul 2020.



Yura Malitsky and Konstantin Mishchenko.

Adaptive proximal gradient method for convex optimization, 2024.



Jorge Nocedal.

Updating quasi-Newton matrices with limited storage.

Mathematics of Computation, 35(151):773–782, 1980.



Peter Ochs, Yunjin Chen, Thomas Brox, and Thomas Pock.

iPiano: Inertial proximal algorithm for nonconvex optimization.

SIAM Journal on Imaging Sciences, 7(2):1388–1419, 2014.



Panagiotis Patrinos and Alberto Bemporad.

Proximal Newton methods for convex composite optimization.

In *52nd IEEE Conference on Decision and Control*, pages 2358–2363, Florence, Italy, 2013.

References V



Pieter Pas, Mathijs Schuurmans, and Panagiotis Patrinos.

Alpaqa: A matrix-free solver for nonlinear MPC and large-scale nonconvex optimization.
In *2022 European Control Conference (ECC)*, pages 417–422, 2022.



Pantelis Sopasakis, Emil Fresk, and Panagiotis Patrinos.

OpEn: Code generation for embedded nonconvex optimization.
IFAC-PapersOnLine, 53(2):6548–6554, 2020.
21st IFAC World Congress.



Lorenzo Stella, Andreas Themelis, Pantelis Sopasakis, and Panagiotis Patrinos.

A simple and efficient algorithm for nonlinear model predictive control.
In *2017 IEEE 56th Annual Conference on Decision and Control (CDC)*, pages 1939–1944, 2017.



Andreas Themelis, Lorenzo Stella, and Panagiotis Patrinos.

Forward-backward envelope for the sum of two nonconvex functions: Further properties and nonmonotone linesearch algorithms.
SIAM Journal on Optimization, 28(3):2274–2303, 2018.



Stephen J. Wright, Robert D. Nowak, and Mario A. T. Figueiredo.

Sparse reconstruction by separable approximation.
IEEE Transactions on Signal Processing, 57(7):2479–2493, 2009.

References VI



Hongchao Zhang and William W. Hager.

A nonmonotone line search technique and its application to unconstrained optimization.

SIAM Journal on Optimization, 14(4):1043–1056, 2004.